

*Regular Articles***Transfer learning with simulated variants and calculated quantities for nanobody melting-temperature prediction**Taihei Murakami^{2,3*}, Kentaro Sasaki^{2*}, Soichiro Oda², Kazuma Okada², Yasuhiro Matsunaga^{1,2}¹ RIKEN Center for Computational Science, Kobe, Hyogo 650-0047, Japan² Saitama University, Saitama 338-8570, Japan³ Epsilon Molecular Engineering, Inc., Saitama 338-8570, Japan

Received XX Xxx, 20XX; Accepted XX Xxx, 20XX;

Released online in J-STAGE as advance publication XX Xxx, 20XX

Edited by Xxx Xxx

Experimental nanobody melting-temperature (T_m) data are limited. Computed stability values can be used during model training, but they describe quantities related to T_m rather than T_m itself. We compared complete simulation plans and calculated quantities as additional training targets. We trained multi-task models that used a shared ESM2 sequence model for experimental T_m and one computed label at a time. We used 57 T_m labels for training, 114 for validation, and 396 reserved for final testing. Each label and encoder setting was tuned separately using only the T_m validation set, then compared using the same test T_m labels. With a frozen encoder, an MD native-contact label lowered T_m error by 0.20°C under a local mutation-scan plan on two fixed structures, compared with an observed reduction of 0.05°C under a heterogeneous structure-panel plan. These plans also differed in contact selection and trajectory averaging, so the comparison does not isolate the choice of variants. Free-energy perturbation (FEP) labels gave the lowest observed error with both frozen and fine-tuned encoders, lowering the fine-tuned mean absolute error from 6.55°C to 6.40°C ; the 95% interval for this change included zero. Native-contact labels gave a clear gain only with the frozen encoder, while Rosetta and ThermoMPNN scores gave little or no gain. The results suggest that simulations for experimental prediction should be planned around both the variants to calculate and the quantity to report, rather than the number of computed values alone.

Key words: nanobody thermostability, melting-temperature prediction, multi-task transfer learning, protein language model, molecular simulation labels

◀ Significance ▶

When experimental melting-temperature data are scarce, computed stability values may help a prediction model. In this study, a local mutation-scan plan gave a larger observed reduction in frozen-encoder error than a heterogeneous structure-panel plan. FEP gave the lowest observed error with both frozen and fine-tuned models. These observations suggest considering both which variants to calculate and which physical quantity to report when planning simulations.

Introduction

Machine-learning models for molecular design need measured examples, but many experiments are slow or expensive. Properties calculated with quantum chemistry, molecular simulation, or other physical models can be used as additional training data. This approach is often called simulation-to-real (Sim2Real) learning. It has been used in materials science

*These authors contributed equally.

Corresponding author: Yasuhiro Matsunaga, RIKEN Center for Computational Science, Kobe, Hyogo 650-0047, Japan.
ORCID iD: <https://orcid.org/0000-0003-2872-3908>, e-mail: ymatsunaga@riken.jp

to combine large computed data sets with smaller experimental data sets [1–3]. In some cases, prediction improves as more computed data are added [3]. This outcome is not guaranteed. A calculation may use an imperfect model, sample the wrong systems, or measure a property that is only weakly related to the experiment.

Protein stability makes this problem clear because several common labels have different physical meanings. Melting temperature (T_m) is the experimental temperature at which folded and unfolded populations are balanced under a given assay. A mutation free-energy change ($\Delta\Delta G$) instead describes the change in folding free energy caused by a sequence change. Rosetta and learned stability models provide scores on their own scales, while molecular dynamics (MD) can provide structural measurements such as the fraction of native contacts that remain during a trajectory. All of these labels concern protein stability, but none is a T_m measurement. A poorly matched computed label may therefore leave T_m prediction unchanged or make it worse.

We study this question in nanobodies, single-domain VHH antibodies used as small and modular binding proteins [4,5]. Thermal stability affects their expression, storage, and formulation [6–8]. Direct measurements and detailed simulations of stability remain costly [9–12]. Protein language models such as ESM2 can learn useful sequence features from large protein databases [13–15], and several methods use these features to predict T_m [16–21]. However, a sequence model does not remove the need for measured T_m values from the protein family of interest.

Large protein-stability data sets are becoming available. For example, Tsuboyama et al. measured folding stability for about 776,000 natural and designed protein domains by cDNA-display proteolysis [22]. Those domains are 40–72 residues long and differ from VHH domains in both sequence and measurement type. Models trained on small domains also transfer less well to longer proteins [23]. Nanobody-specific data and validation are therefore still needed [24–26].

The open question is not simply whether computed labels can be added to a T_m model. It is how to plan the simulations that produce those labels. A mixed set of unrelated sequences may contain unwanted differences such as sequence length. A mutation scan on one or two fixed structures gives a more controlled set of sequence changes. The calculated quantity also matters: a mutation free energy, a Rosetta score, and an MD native-contact value may differ in how closely they relate to T_m . Both the systems to calculate and the reported quantity are chosen before model training.

Here, we compare several computed labels in the same low-data T_m task. The labels are mutation free energies from free-energy perturbation (FEP), Rosetta mutation scores, ThermoMPNN scores, Rosetta scores for ESM2-proposed or random variants, and MD native-contact values. For the MD label, we compare two complete simulation plans: a heterogeneous structure panel and local mutation scans on the same 1MEL and 4IDL structures used for FEP. Both use 400 K MD and native-contact labels, but they also differ in contact selection and trajectory averaging. Their comparison therefore does not isolate the choice of variants. Each computed label is trained as a second task that shares an ESM2 sequence model with T_m prediction (Fig. 1). We tune every label separately using only the experimental T_m validation set. We then compare the selected models using paired errors on the same 396 test T_m labels, which were not used for training or model selection. In this paper, “transfer” means that adding a computed label lowers the test T_m error.

We observed a larger frozen-encoder gain under the local mutation-scan MD plan than under the heterogeneous-panel plan. Among the calculated quantities, FEP gave the lowest observed error in both encoder settings, whereas native-contact values gave a clear gain only with the frozen encoder. These results suggest that simulations intended to support an experimental predictor should be planned around both the systems to calculate and the quantity to report.

Materials and methods

Study design

We asked whether computed protein-stability labels could improve nanobody melting-temperature (T_m) prediction when only a small experimental training set was available. Experimental T_m validation error, never computed-label error, was used to select checkpoints and hyperparameters. Each computed label was added as a second regression task that shared all or part of the sequence model with the T_m task. We compared each label with a T_m -only model trained and selected in the same encoder setting.

Experimental T_m data

We used the thermo- t_m data set from NbBench, a public nanobody benchmark [25]. We reassigned its three published partitions before the comparisons in this study so that the smallest partition was used for training. The resulting split contained 57 training, 114 validation, and 396 test sequences. This split models the setting that motivated the study: few T_m measurements for fitting a model, but enough test proteins for paired comparisons. The training set fit model parameters, the validation set selected checkpoints and hyperparameters, and the test set was kept for final evaluation.

T_m values were stored in degrees Celsius. We fit a robust scaler and then a min–max scaler to [0,1] using only the 57




Figure 1 Multi-task model used to test computed stability labels. Experimental sequences and sequences with computed labels pass through the same ESM2 model and shared multilayer perceptron. Separate scalar heads predict T_m and the computed label. The ESM2 model is either kept frozen or fine-tuned. Models are selected using experimental T_m validation data, and the final comparison uses held-out T_m test data.

training values. The fitted scalers were applied to the training and validation labels. Predictions were returned to degrees Celsius before validation and test errors were calculated. Validation and test T_m values therefore did not affect the scaling step.

Computed-label tables and sampling

We tested seven sources of computed labels: FEP mutation free energies, matched MD native-contact values, Rosetta mutation scores, ThermoMPNN predictions, Rosetta scores for random variants, Rosetta scores for ESM2-proposed variants, and MD native-contact values for a heterogeneous structure panel. These values were treated as related measurements, not as estimates of T_m.

FEP, Rosetta, and ThermoMPNN used two single-mutation tables based on the 1MEL and 4IDL structures. Each contained 435 and 409 rows, respectively. The matched MD tables used the same two structures and mutation scans and contained 431 1MEL rows and 406 4IDL rows that reached the common analysis time. All 837 matched-MD sequences were present in the 844-row FEP pool; seven FEP variants did not have a matched MD label. The random and ESM2-proposed variant sets each contained 1000 two-mutation sequences per structure. The heterogeneous MD table contained 1143 PDB-derived rows representing 833 unique amino-acid sequences. The sampling unit for this table was a PDB row, so a sequence represented by more than one structure could occur more than once.

For all two-table sources, the notation n means a cap of n rows from *each* structure table for each ensemble member, not n rows in total. Thus, at $n = 320$, each member sampled at most 640 computed-label rows. The heterogeneous MD source

had one table, so n is the total number sampled for each member. Each run used its own seed to sample rows. Sampling was carried out separately for each label source; in particular, the FEP and matched-MD runs were not forced to select the same mutations from their nearly identical pools. After sampling, 80% of the rows were used for training and 20% were set aside. The set-aside rows were not used to choose checkpoints; only the experimental T_m validation rows were used for that purpose. The relation between sequence length and the heterogeneous-panel MD label is shown in Supplementary Fig. S1; exact sequence matches to the target data are reported below.

The 1MEL and 4IDL tables were kept as two tasks because their score ranges can differ. Models used either two separate output heads or one output head with a learned structure indicator, as described below. Every processed table supplied an amino-acid sequence and one scalar label scaled to [0,1]. FEP, Rosetta, and ThermoMPNN scores were sign-reversed so that larger values indicated a more favorable change, then processed separately for each structure with robust scaling, a Yeo–Johnson transform, and min–max scaling. MD contact values were min–max scaled with larger values indicating greater contact retention.

We also checked exact sequence matches between every computed-label table and the NbBench splits. There were no exact matches for the FEP, matched MD, Rosetta, ThermoMPNN, random-variant, or ESM2-proposed tables. The heterogeneous MD table contained 2 training, 4 validation, and 8 test sequences also present in NbBench. When one of these rows was sampled for auxiliary training, the model saw the target sequence with a computed Q label. It did not see the experimental validation or test T_m value, which remained unused in training. We retained the rows and report them because the heterogeneous-panel comparison was therefore not fully disjoint from the target sequences.

FEP mutation free energies

The FEP labels were mutation free-energy changes, $\Delta\Delta G$, for the 1MEL and 4IDL mutation scans. Positive raw $\Delta\Delta G$ denotes a destabilizing mutation.

The fixed FEP tables were produced with alchemical simulations in NAMD using the CHARMM36 protein force field and explicit TIP3P water [27–29]. The starting coordinates were the 1MEL and 4IDL crystal structures. Mutation systems were prepared with the VMD AlaScan workflow and CHARMM hybrid-residue topologies [30, 31]. Representative archived input files show a 2-fs timestep, particle mesh Ewald electrostatics, bonds to hydrogen held fixed, a temperature of 300 K, and a pressure of 1 atm [32, 33].

Each forward and backward transformation covered 20 windows of width $\Delta\lambda = 0.05$ between 0 and 1. A window contained 500,000 steps (1 ns), with its first 100,000 steps assigned to equilibration. The folded-state free energy was combined with a residue-specific unfolded-state reference to obtain $\Delta\Delta G$ relative to the wild type.

MD native-contact labels

Both MD labels were based on the smooth native-contact fraction Q [34]. For a native contact with reference distance d_{ij}^0 and distance $d_{ij}(t)$ at time t , its contribution was

$$q_{ij}(t) = \left[1 + \exp \left\{ \beta [d_{ij}(t) - \lambda d_{ij}^0] \right\} \right]^{-1},$$

with $\beta = 50 \text{ nm}^{-1}$ and $\lambda = 1.8$. Native contacts were pairs within 0.45 nm in the reference frame and separated by more than three residues. Q was the mean over the selected contacts and frames.

Matched mutation scans. We simulated the 1MEL and 4IDL wild types and the same single-mutant sequences used in the FEP comparison. Each protein was placed in a cubic box of OPC water with a 15-Å buffer and 0.15 M salt, and was described by Amber ff19SB with hydrogen-mass repartitioning [35, 36]. OpenMM was used for minimization, equilibration, and production [37]. The recorded run settings were 0.2 ns NVT followed by 0.8 ns NPT at 300 K, 0.5 ns restrained NVT at 400 K, and restraint-free NVT production at 400 K. All stages used a 4-fs timestep; production coordinates were written every 10 ps. Although some runs continued to 100 ns, the label used the first 4000 production frames (40 ns) for every variant.

For this source, the reference distances came from production frame 0 and all backbone atoms selected by MDAnalysis were eligible for contacts. We computed $\Delta Q = Q_{\text{mutant}} - Q_{\text{WT}}$ within each structure and then min–max scaled it. Subtracting one constant wild-type value does not change a min–max-scaled label, so scaled Q and scaled ΔQ are identical within each structure table.

Heterogeneous nanobody panel. The comparison panel was built from SABDab structures [38]. One verified heavy-chain nanobody was taken per PDB entry when an IMGT-renumbered structure was available. Preparation used PDBFixer and pH-aware protonation with pdb2pqr/PROPKA when available [39, 40]. The systems used ff19SB, OPC water, 0.15

M salt, and OpenMM. After 300 K NVT and NPT equilibration, two restrained 400 K NVT stages preceded 100 ns of unrestrained 400 K NVT production. Requiring a usable topology and trajectory, a sequence of at least 50 residues, and at least one selected contact gave 1143 rows.

The heterogeneous-panel label used production frame 0 as its reference and averaged approximately the last 30 ns. Its contact set contained backbone heavy-atom pairs for which at least one residue was Asp, Glu, Gln, Asn, Arg, Lys, or His. It was min–max scaled across the single table. The matched and heterogeneous studies used trajectory lengths in the same broad range: tens of nanoseconds for the matched scans and 100 ns for the heterogeneous panel. They used the same Best–Hummer form and the same temperature, although their contact selection and averaging window were not identical. We therefore treat this as a comparison of complete simulation plans, rather than as a one-variable test of sequence choice alone.

Rosetta, ThermoMPNN, and variant-selection comparisons

Rosetta scores were calculated with the `ddg_monomer` application and the `ref2015` score function [41,42]. The archived command used 50 iterations, Cartesian minimization, minimum-score aggregation, and full-atom input. The same settings were used for the 1MEL/4IDL mutation tables and for the two sets of two-mutation variants.

For the random set, mutation positions and replacement amino acids were sampled uniformly, excluding cysteine as a replacement, until 1000 unique two-mutation variants had been obtained for each structure. For the ESM2-guided set, the 650M-parameter ESM2 model generated 100,000 two-mutation variants per structure [14]. Variants were ranked by pseudo-log-likelihood, and the top 1% were scored with Rosetta. ThermoMPNN predictions were made for the same 435 and 409 single-mutation sequences used by the structure-matched tables [43]. We used the two processed ThermoMPNN tables, not the larger raw Colab outputs stored with them.

Sequence model and regression heads

The main encoder was the 8M-parameter ESM2 model `facebook/esm2_t6_8M_UR50D` [14]. The tokenizer used a maximum length of 160 tokens, including the beginning and end tokens, so at most 158 amino acids were supplied to ESM2. Longer sequences were truncated at the end. All NbBench target sequences were 151 residues or shorter. The first-token representation was used as the sequence vector. In the shared model, this vector passed through linear layers of size 256, 128, and 32, with ReLU activation and dropout after the first two layers. The Tm head and each computed-label head were separate 32-to-1 linear layers.

The model search also included residual and latent forms. Both gave Tm and computed-label examples separate 256–128–32 paths. The residual form added a sigmoid-gated correction to a base Tm prediction. The latent form predicted a Tm correction from the Tm vector, the second path’s vector, and their elementwise product. For two-table sources, we compared separate 1MEL and 4IDL heads with a single head conditioned on a learned structure embedding. In the residual and latent implementations, the computed-label loss received a detached encoder vector by default, whereas the Tm path could still update the encoder. In the shared form, losses from all tasks could update the encoder when it was not frozen.

We tested two encoder settings. In the frozen setting, all ESM2 parameters were fixed and sequence vectors were computed once before head training. In the fine-tuned setting, all ESM2 parameters could change. AdamW used a smaller learning rate for ESM2 than for the newly initialized layers. Selected controls used 35M- and 650M-parameter ESM2 encoders.

Loss and model selection

Each row had one task identifier and contributed loss only through that task’s output head. We used Huber loss with $\delta = 1$ on the scaled labels. By default, the losses were combined with learned uncertainty weights [44]. For task k , loss L_k , and learned log-scale s_k , the term added to the total loss was

$$\frac{L_k}{2 \exp(2s_k)} + s_k.$$

Fixed task weights were tested as controls.

Models were trained with AdamW, cosine learning-rate decay, a batch size of 16, 100 warmup steps, and at most 400 epochs. Mixed-precision training was used on CUDA devices. A checkpoint was saved and evaluated once per epoch. The checkpoint with the lowest mean squared error on the 114 experimental Tm validation rows was retained, and training stopped after 15 epochs without improvement.

For the structure-matched label comparison in Fig. 3, each label source was tuned separately in the frozen and fine-tuned

settings. The search had two stages. Stage 1 crossed three model forms (shared, residual, and latent) with two head forms (separate and structure-conditioned); Tm-only models searched the three model forms without a computed-label head. Stage 2 held the chosen form fixed. For fine-tuned models it tested encoder learning rates of 1×10^{-4} , 5×10^{-5} , 3×10^{-5} , and 1×10^{-5} . For frozen models it tested head learning rates of 1×10^{-3} , 5×10^{-4} , and 2×10^{-4} . Both searches crossed dropout values of 0.05, 0.15, and 0.30 with weight decay values of 0.02 and 0.10.

Each candidate was ranked by the MAE of a three-seed ensemble on the experimental Tm validation set, using $n = 320$ rows per structure table. The selected setting was then trained with five seeds for the final test result. Label-count curves used the same selected setting at $n = 20, 80, 160$, and 320 per structure; they were not retuned at each point. The heterogeneous-panel study was selected separately with its own Tm-only reference. Its final ensembles used ten seeds and $n = 640$ rows from the single heterogeneous table, the same total number sampled per ensemble member as $n = 320$ from each of the two matched tables. The run seed controlled randomness in Python, NumPy, PyTorch, and CUDA, as well as computed-row sampling and the computed-label training–monitoring split.

Test evaluation and statistics

For each condition, the five matched-label models, or the ten heterogeneous-panel models, predicted all 396 test sequences. Their predictions were averaged in scaled space and then returned to degrees Celsius. The main measure was mean absolute error (MAE). We also calculated root mean squared error, R^2 , and Spearman and Pearson correlations.

Uncertainty was estimated by nonparametric bootstrap over test proteins. For one condition, we resampled its 396 absolute errors with replacement 10,000 times, using random seed 42, and calculated MAE for each sample. The reported 95% confidence interval is the 2.5th to 97.5th percentile range.

For a comparison, the same resampled test indices were applied to the candidate and reference error vectors. We calculated Δ MAE as candidate MAE minus reference MAE for every sample. Negative values therefore favor the candidate. Main figures report the mean paired Δ MAE and its 95% confidence interval. We do not use a separate p-value threshold for the main claims.

Software and saved outputs

The data loader, training code, model settings, and ensemble evaluation are in `prepare.py`, `train.py`, and the scripts distributed with the repository. Each final run stores its command-line settings, predictions, absolute errors, summary measures, and bootstrap results in machine-readable files. The figure scripts read those saved files directly. The released procedure begins with the fixed processed CSV tables; it does not repeat the raw FEP, Rosetta, ThermoMPNN, or MD calculations.

Results and discussion

Computed labels were judged only by experimental Tm

We tested each computed label as a second regression task (Fig. 1). The computed values were not treated as extra Tm measurements. ESM2 was either frozen or fine-tuned, and only the Tm validation error was used to choose models and settings. We used 57 measured Tm values for training, 114 for validation, and 396 for final testing. The paired comparisons used the same Tm test set for every model.

The selected Tm-only model had a test MAE of 7.23°C with a frozen encoder and 6.55°C with a fine-tuned encoder. We used each baseline only for comparisons within the same encoder setting. We then asked two questions. First, how did the results differ between the two MD plans? Second, which computed labels gave the most consistent results when ESM2 was frozen or fine-tuned?

The matched mutation-scan plan gave a larger frozen-encoder gain

We compared two MD simulation plans. The heterogeneous plan sampled 1143 PDB structures representing 833 unique sequences of different lengths. The matched plan used local mutations of the fixed 1MEL and 4IDL structures. Both used 400 K MD, trajectory lengths of tens to 100 ns per sequence, and labels from the Best–Hummer native-contact Q family. Their contact definitions, averaging windows, and model searches were not identical. This is therefore a comparison of two complete plans, rather than a test in which only the selected sequences were changed.

The two studies were run and tuned separately, so we did not compare their absolute MAEs. Figure 2a instead shows the paired change from the Tm-only model in each study. In the heterogeneous study, the MD label changed MAE by -0.049°C with a frozen encoder (95% CI, -0.092 to -0.006°C) and by $+0.116^\circ\text{C}$ with a fine-tuned encoder (95% CI, -0.009 to $+0.243^\circ\text{C}$). In the matched study, the frozen change was -0.195°C (95% CI, -0.366 to -0.030°C), a larger




Figure 2 The matched mutation-scan plan showed a larger observed frozen-encoder gain from native-contact Q . The heterogeneous plan distributed simulations across PDB structures of different lengths, whereas the matched plan used local mutations of 1MEL and 4IDL. Both used 400 K MD and labels from the same Q family, but their contact definitions, averaging windows, model searches, and Tm-only references differed. (a) Paired MAE change for each plan and encoder setting. Zero means equal MAE to that study's own Tm-only model, and negative values mean lower Tm error. The heterogeneous and matched studies used fixed 10-seed and five-seed ensembles, respectively. (b) Paired MAE change from the frozen Tm-only model over four label counts. Here n is the number sampled from each structure table for each ensemble member; 80% of the sampled rows entered training. Points and whiskers show the paired change and its 95% bootstrap interval over the 396 Tm test values.

observed gain than in the heterogeneous study. The fine-tuned change was $+0.029^{\circ}\text{C}$ (95% CI, -0.139 to $+0.197^{\circ}\text{C}$).

The heterogeneous table contained eight sequences from the Tm test set. The model saw computed Q values, but not the test Tm values, for these sequences. When the eight sequences were omitted from the error calculation, the paired changes were almost unchanged: -0.049°C for the frozen encoder (95% CI, -0.094 to -0.005°C) and $+0.122^{\circ}\text{C}$ for the fine-tuned encoder (95% CI, -0.005 to $+0.248^{\circ}\text{C}$). Thus, the test errors of these eight sequences did not drive the contrast in Fig. 2a. We did not retrain after removing the overlapping computed-label rows, however, so this check does not rule out an effect on training or model selection.

Several features of the two plans may have contributed. In the heterogeneous table, Q had a Pearson correlation of $r = +0.13$ with sequence length (Supplementary Fig. S1), so the model could partly use length instead of the intended contact information. The matched scans held length fixed within each structure and described local changes caused by one mutation. The heterogeneous table also sampled PDB rows, so repeated structures of one sequence received more weight. These observations do not show that any one feature caused the difference.

Subtracting a structure-specific wild-type constant does not affect later min-max scaling, so scaled Q and scaled ΔQ were identical within each matched table.

The frozen label-count curves provide a second view of the matched plan (Fig. 2b). At $n = 20$, the paired changes were -0.10°C for FEP and $+0.09^{\circ}\text{C}$ for matched MD. At $n = 320$, they were -0.22°C for FEP and -0.20°C for matched MD, and both 95% intervals were below zero. The direct FEP-minus-MD difference at $n = 320$ was -0.03°C (95% CI, -0.24 to $+0.19^{\circ}\text{C}$), so we did not resolve a difference between the two labels with a frozen encoder. The four settings were not monotonic enough to support a scaling law.

FEP gave the lowest observed error in both encoder settings

We next compared all computed labels after each was tuned separately (Fig. 3a). With a frozen encoder, FEP and matched MD gave the clearest gains. FEP reached 7.01°C MAE, a paired change of -0.221°C from Tm-only (95% CI, -0.393 to -0.051°C). Matched MD reached 7.03°C, a change of -0.195°C (95% CI, -0.366 to -0.030°C). ThermoMPNN was third at 7.09°C, but its interval included zero. The Rosetta-based labels were close to or worse than the Tm-only model.

With a fine-tuned encoder, FEP reached 6.40°C, the lowest observed MAE in the study. Its paired change from the 6.55°C Tm-only model was -0.153°C (95% CI, -0.323 to $+0.020^{\circ}\text{C}$). The interval included zero. Matched MD, ThermoMPNN, plain Rosetta, and random variants plus Rosetta differed from Tm-only by $+0.029$ to $+0.144^{\circ}\text{C}$, with intervals that also included zero. ESM2-proposed variants plus Rosetta had higher error by $+0.411^{\circ}\text{C}$ (95% CI, $+0.140$ to $+0.698^{\circ}\text{C}$).

Figure 3b directly compares FEP and matched MD. Zero means equal Tm test-set MAE, and negative values favor FEP. The frozen difference was -0.026°C (95% CI, -0.243 to $+0.192^{\circ}\text{C}$). With a fine-tuned encoder, the observed difference was -0.182°C (95% CI, -0.393 to $+0.024^{\circ}\text{C}$). Neither interval excluded zero at the largest label count.

The label-count curves show how these observed differences developed (Fig. 3c). At $n = 20$, matched MD had higher error than Tm-only by $+0.47^{\circ}\text{C}$ (95% CI, $+0.26$ to $+0.69^{\circ}\text{C}$). Its error decreased as n increased, and its interval included zero by $n = 160$. FEP changed from $+0.02^{\circ}\text{C}$ at $n = 20$ to -0.15°C at $n = 320$, but the 95% interval at the largest setting also included zero. These four settings show recovery and crossover, not a power law.

FEP $\Delta\Delta G$ and Tm are different quantities, but both are linked to folding thermodynamics. Here, Tm supplied absolute temperatures across nanobodies, whereas FEP supplied local stabilizing and destabilizing changes in two structures. Native-contact retention is a more indirect measure of folding stability. This difference offers one possible explanation for the frozen and fine-tuned results, but the present comparison cannot establish the mechanism.

A recent preprint from our group examined ESM2 after fine-tuning on nanobody Tm with sparse autoencoders [45]. It found features linked to surface charge, hydrophobic-core packing, the VHH tetrad, and disulfide bonds. Together with that work, the present result suggests a testable question: do FEP labels support some of the same local stability features that appear during Tm fine-tuning? We did not examine hidden representations here, and the selected FEP and MD models did not always use the same regression form. Comparing the learned features of Tm-only, Tm+FEP, and Tm+MD models is therefore a separate next step.

Checks and limits

Exploratory fixed-setting tests with the 35M- and 650M-parameter ESM2 models showed the same direction as the selected 8M model: adding FEP gave a lower observed MAE than Tm-only within each model size (Supplementary Fig. S2a). The larger models were not searched in the same way as the 8M model, so these tests do not provide a model-size comparison. An analytical correction for periodic images in charge-changing FEP mutations had almost no effect on the scaled labels (Supplementary Fig. S2b).

ESM2-proposed variants scored by Rosetta did not have lower error than random variants scored by the same method. This result does not test a full protein design study; it shows that proposing a sequence and scoring it must be judged together.

Computed data should be judged by how well they help the measured task, and the choice of proteins and calculations should be made before spending the simulation time. A large mixed set is not automatically better than a smaller local scan. FEP and matched MD used roughly similar simulation time per variant here, but we did not save wall-clock costs in a form that supports a broader cost claim.

This study has several limits. The Tm training set contains only 57 nanobodies, and the matched labels cover two structures. The two MD plans differ in more than their sequence sets. The heterogeneous table sampled repeated PDB entries, contained eight exact matches to the Tm test sequences, and included 11 unusually long rows. The selected FEP and MD models also differed in regression form and in whether the auxiliary loss updated the encoder. The bootstrap intervals resample the fixed set of test proteins after averaging the trained models; they do not include uncertainty from choosing auxiliary rows, training seeds, or model settings. Finally, we did not make new Tm measurements or test predictions on newly measured nanobodies.




Figure 3 FEP gave the lowest observed Tm error in both encoder settings. (a) Paired MAE change from the corresponding Tm-only model for every tuned label. Negative values mean lower Tm error. Filled squares and open circles denote frozen and fine-tuned encoders, respectively. (b) Direct FEP-minus-MD comparison. Zero means equal Tm test-set MAE, negative values favor FEP, and positive values favor MD native-contact Q . (c) Paired MAE change from the fine-tuned Tm-only model over four label counts. Here n is the number sampled from each structure table for each ensemble member; 80% entered training. Whiskers in all panels are 95% bootstrap intervals over the 396 Tm test values.

Conclusion

Computed labels improved low-data nanobody T_m prediction only in specific settings. Among the complete MD plans tested here, native-contact labels from local mutation scans on 1MEL and 4IDL gave a larger observed frozen-encoder gain than labels from a heterogeneous structure panel. The plans differed in more than the choice of variants, so this result does not assign the difference to one factor. FEP mutation free energies gave the lowest observed error in both encoder settings, although the 95% interval for the fine-tuned comparison included zero. Matched-scan MD native-contact values gave a clear gain only with the frozen model. Together, these results suggest planning simulations around both the systems to calculate and the quantity to report, rather than around data volume alone.

In the longer term, matched scans could be extended to more structures and used in an iterative design study. A model could choose which variant to calculate next, receive an FEP or MD result, and use that result to choose the following variant. Reinforcement learning is one way to make these sequential choices, but the present results caution against using a computed score as the only reward. Our comparison of ESM2 and random proposals under the same Rosetta score shows why variant generation and scoring must be tested together. Any future reinforcement-learning study should therefore be compared with random sampling and simpler active-learning methods, and its success should be tested with new experimental T_m measurements.

Conflict of interest

T.M. is an employee of Epsilon Molecular Engineering, Inc. Y.M. is a scientific advisor of Epsilon Molecular Engineering, Inc.

Author contributions

Y.M. conceived and designed the study. Y.M., T.M., and K.S. carried out the overall computational study. Y.M. and T.M. analyzed the data. S.O. and K.O. performed the FEP calculations. Y.M. and T.M. wrote the manuscript. All authors discussed the results, reviewed the manuscript, and approved the final version.

Data availability

The source code for data processing, model training, analysis, and figure generation is available at <https://github.com/matsunagalab/sim2real>. The processed label tables, per-sequence test errors, and figure source data will be available from Zenodo (DOI: 10.5281/zenodo.XXXXXXX). Raw MD and FEP files are not included because of their size and are available from the corresponding author upon reasonable request.

Acknowledgements

This work was supported by JSPS KAKENHI Grant 23H03412 and JST NBDC Grant JPMJND2603, and partly supported by MEXT as “Program for Promoting Researches on the Supercomputer Fugaku” (Development and application of large-scale simulation-based inferences for biomolecules JPMXP1020230119). Computational resources were provided by the supercomputer Fugaku at the RIKEN Center for Computational Science under Project IDs hp230209, hp240215, and hp250233.

References

- [1] Yamada, H., Liu, C., Wu, S., Koyama, Y., Ju, S., Shiomi, J., et al. Predicting materials properties with little data using shotgun transfer learning. *ACS Central Science* 5, 1717–1730 (2019). <https://doi.org/10.1021/acscentsci.9b00804>
- [2] Aoki, Y., Wu, S., Tsurimoto, T., Hayashi, Y., Minami, S., Okubo, T., et al. Multitask machine learning to predict polymer–solvent miscibility using Flory–Huggins interaction parameters. *Macromolecules* 56, 5446–5456 (2023). <https://doi.org/10.1021/acs.macromol.2c02600>
- [3] Minami, S., Hayashi, Y., Wu, S., Fukumizu, K., Sugisawa, H., Ishii, M., et al. Scaling law of Sim2Real transfer learning in expanding computational materials databases for real-world predictions. *npj Computational Materials* 11, 146 (2025). <https://doi.org/10.1038/s41524-025-01606-5>
- [4] Muyldermans, S. Nanobodies: Natural single-domain antibodies. *Annual Review of Biochemistry* 82, 775–797 (2013). <https://doi.org/10.1146/annurev-biochem-063011-092449>

- [5] Alexander, E. Leong, K. W. Discovery of nanobodies: a comprehensive review of their applications and potential over the past five years. *Journal of Nanobiotechnology* 22, 661 (2024). <https://doi.org/10.1186/s12951-024-02900-y>
- [6] Kunz, P., Zinner, K., Mücke, N., Bartoschik, T., Muyldermans, S., Hoheisel, J. D. The structural basis of nanobody unfolding reversibility and thermoresistance. *Sci. Rep.* 8, 7934 (2018).
- [7] Mezzasalma, T. M., Kranz, J. K., Chan, W., Struble, G. T., Schalk-Hihi, C., Deckman, I. C., et al. Enhancing recombinant protein quality and yield by protein stability profiling. *J. Biomol. Screen.* 12, 418–428 (2007).
- [8] Kumar, S., Tsai, C. J., Nussinov, R. Factors enhancing protein thermostability. *Protein Eng.* 13, 179–191 (2000).
- [9] Bekker, G.-J., Ma, B., Kamiya, N. Thermal stability of single-domain antibodies estimated by molecular dynamics simulations. *Protein Science* 28, 429–438 (2019). <https://doi.org/10.1002/pro.3546>
- [10] Higashida, R. Matsunaga, Y. Enhanced conformational sampling of nanobody CDR H3 loop by generalized replica-exchange with solute tempering. *Life (Basel)* 11, 1428 (2021).
- [11] Mohseni, A., Molakarimi, M., Taghdir, M., Sajedi, R. H., Hasannia, S. Exploring single-domain antibody thermostability by molecular dynamics simulation. *J. Biomol. Struct. Dyn.* 37, 3686–3696 (2019).
- [12] Tomimoto, Y., Yamazaki, R., Shirai, H. Increasing the melting temperature of VHH with the in silico free energy score. *Sci. Rep.* 13, 4922 (2023).
- [13] Rives, A., Meier, J., Sercu, T., Goyal, S., Lin, Z., Liu, J., et al. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences of the United States of America* 118, e2016239118 (2021). <https://doi.org/10.1073/pnas.2016239118>
- [14] Lin, Z., Akin, H., Rao, R., Hie, B., Zhu, Z., Lu, W., et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science* 379, 1123–1130 (2023). <https://doi.org/10.1126/science.ade2574>
- [15] Hayes, T., Rao, R., Akin, H., Sofroniew, N. J., Oktay, D., Lin, Z., et al. Simulating 500 million years of evolution with a language model. *Science* 387, 850–858 (2025).
- [16] Blaabjerg, L. M., Kassem, M. M., Good, L. L., Jonsson, N., Cagiada, M., Johansson, K. E., et al. Rapid protein stability prediction using deep learning representations. *eLife* 12, e82593 (2023). <https://doi.org/10.7554/eLife.82593>
- [17] Li, M., Wang, H., Yang, Z., Zhang, L., Zhu, Y. DeepTM: A deep learning algorithm for prediction of melting temperature of thermophilic proteins directly from sequences. *Computational and Structural Biotechnology Journal* 21, 5544–5560 (2023). <https://doi.org/10.1016/j.csbj.2023.11.006>
- [18] Alvarez, J. A. E. Dean, S. N. TEMPRO: nanobody melting temperature estimation model using protein embeddings. *Scientific Reports* 14, 19074 (2024). <https://doi.org/10.1038/s41598-024-70101-6>
- [19] Savojardo, C., Manfredi, M., Martelli, P. L., Casadio, R. DDGemb: predicting protein stability change upon single- and multi-point variations with embeddings and deep learning. *Bioinformatics* 41 (2024).
- [20] Ramon, A., Ni, M., Predeina, O., Gaffey, R., Kunz, P., Onuoha, S., et al. Prediction of protein biophysical traits from limited data: a case study on nanobody thermostability through NanoMelt. *MAbs* 17, 2442750 (2025).
- [21] Mao, J., Song, Y., Kong, M., Guo, Y., Liu, Y., Pu, X. Thermostability prediction powered by synergistic deep learning at experimental and theoretical levels for nanobodies. *ACS Appl. Mater. Interfaces* 18, 6433–6451 (2026).
- [22] Tsuboyama, K., Dauparas, J., Chen, J., Laine, E., Mohseni Behbahani, Y., Weinstein, J. J., et al. Mega-scale experimental analysis of protein folding stability in biology and design. *Nature* 620, 434–444 (2023). <https://doi.org/10.1038/s41586-023-06328-6>
- [23] Chu, S. K. S., Narang, K., Siegel, J. B. Protein stability prediction by fine-tuning a protein language model on a mega-scale dataset. *PLOS Computational Biology* 20, e1012248 (2024). <https://doi.org/10.1371/journal.pcbi.1012248>
- [24] Valdés-Tresanco, M. S., Valdés-Tresanco, M. E., Molina-Abad, E., Moreno, E. NbThermo: a new thermostability database for nanobodies. *Database* 2023, baad021 (2023). <https://doi.org/10.1093/database/baad021>
- [25] Zhang, Y. Tsuda, K. NbBench: Benchmarking language models for comprehensive nanobody tasks. *arXiv* 2505.02022 (2025). <https://doi.org/10.48550/arXiv.2505.02022>
- [26] Schmirler, R., Heinzinger, M., Rost, B. Fine-tuning protein language models boosts predictions across diverse tasks. *Nature Communications* 15, 7407 (2024). <https://doi.org/10.1038/s41467-024-51844-2>
- [27] Phillips, J. C., Hardy, D. J., Maia, J. D. C., Stone, J. E., Ribeiro, J. V., Bernardi, R. C., et al. Scalable molecular dynamics on CPU and GPU architectures with NAMD. *The Journal of Chemical Physics* 153, 044130 (2020). <https://doi.org/10.1063/5.0014475>
- [28] Best, R. B., Zhu, X., Shim, J., Lopes, P. E. M., Mittal, J., Feig, M., et al. Optimization of the additive CHARMM all-atom protein force field targeting improved sampling of the backbone ϕ , ψ and side-chain χ_1 and χ_2 dihedral angles. *Journal of Chemical Theory and Computation* 8, 3257–3273 (2012). <https://doi.org/10.1021/ct300400x>
- [29] Jorgensen, W. L., Chandrasekhar, J., Madura, J. D., Impey, R. W., Klein, M. L. Comparison of simple potential functions for simulating liquid water. *The Journal of Chemical Physics* 79, 926–935 (1983).

- <https://doi.org/10.1063/1.445869>
- [30] Humphrey, W., Dalke, A., Schulten, K. VMD: Visual molecular dynamics. *Journal of Molecular Graphics* 14, 33–38 (1996). [https://doi.org/10.1016/0263-7855\(96\)00018-5](https://doi.org/10.1016/0263-7855(96)00018-5)
 - [31] Ramadoss, V., Dehez, F., Chipot, C. AlaScan: A graphical user interface for alanine scanning free-energy calculations. *Journal of Chemical Information and Modeling* 56, 1122–1126 (2016). <https://doi.org/10.1021/acs.jcim.6b00162>
 - [32] Essmann, U., Perera, L., Berkowitz, M. L., Darden, T., Lee, H., Pedersen, L. G. A smooth particle mesh Ewald method. *The Journal of Chemical Physics* 103, 8577–8593 (1995). <https://doi.org/10.1063/1.470117>
 - [33] Ryckaert, J.-P., Ciccotti, G., Berendsen, H. J. C. Numerical integration of the cartesian equations of motion of a system with constraints: molecular dynamics of n-alkanes. *Journal of Computational Physics* 23, 327–341 (1977). [https://doi.org/10.1016/0021-9991\(77\)90098-5](https://doi.org/10.1016/0021-9991(77)90098-5)
 - [34] Best, R. B., Hummer, G., Eaton, W. A. Native contacts determine protein folding mechanisms in atomistic simulations. *Proceedings of the National Academy of Sciences of the United States of America* 110, 17874–17879 (2013). <https://doi.org/10.1073/pnas.1311599110>
 - [35] Tian, C., Kasavajhala, K., Belfon, K. A. A., Raguette, L., Huang, H., Miguez, A. N., et al. ff19SB: amino-acid-specific protein backbone parameters trained against quantum mechanics energy surfaces in solution. *Journal of Chemical Theory and Computation* 16, 528–552 (2020). <https://doi.org/10.1021/acs.jctc.9b00591>
 - [36] Izadi, S., Anandakrishnan, R., Onufriev, A. V. Building water models: a different approach. *The Journal of Physical Chemistry Letters* 5, 3863–3871 (2014). <https://doi.org/10.1021/jz501780a>
 - [37] Eastman, P., Swails, J., Chodera, J. D., McGibbon, R. T., Zhao, Y., Beauchamp, K. A., et al. OpenMM 7: rapid development of high performance algorithms for molecular dynamics. *PLOS Computational Biology* 13, e1005659 (2017). <https://doi.org/10.1371/journal.pcbi.1005659>
 - [38] Dunbar, J., Krawczyk, K., Leem, J., Baker, T., Fuchs, A., Georges, G., et al. SABDab: the structural antibody database. *Nucleic Acids Research* 42, D1140–D1146 (2014). <https://doi.org/10.1093/nar/gkt1043>
 - [39] Dolinsky, T. J., Nielsen, J. E., McCammon, J. A., Baker, N. A. PDB2PQR: an automated pipeline for the setup of Poisson–Boltzmann electrostatics calculations. *Nucleic Acids Research* 32, W665–W667 (2004). <https://doi.org/10.1093/nar/gkh381>
 - [40] Olsson, M. H. M., Søndergaard, C. R., Rostkowski, M., Jensen, J. H. PROPKA3: consistent treatment of internal and surface residues in empirical pK_a predictions. *Journal of Chemical Theory and Computation* 7, 525–537 (2011). <https://doi.org/10.1021/ct100578z>
 - [41] Kellogg, E. H., Leaver-Fay, A., Baker, D. Role of conformational sampling in computing mutation-induced changes in protein structure and stability. *Proteins: Structure, Function, and Bioinformatics* 79, 830–838 (2011). <https://doi.org/10.1002/prot.22921>
 - [42] Alford, R. F., Leaver-Fay, A., Jeliazkov, J. R., O’Meara, M. J., DiMaio, F. P., Park, H., et al. The Rosetta all-atom energy function for macromolecular modeling and design. *Journal of Chemical Theory and Computation* 13, 3031–3048 (2017). <https://doi.org/10.1021/acs.jctc.7b00125>
 - [43] Dieckhaus, H., Brocidiaco, M., Randolph, N. Z., Kuhlman, B. Transfer learning to leverage larger datasets for improved prediction of protein stability changes. *Proceedings of the National Academy of Sciences of the United States of America* 121, e2314853121 (2024). <https://doi.org/10.1073/pnas.2314853121>
 - [44] Kendall, A., Gal, Y., Cipolla, R. Multi-task learning using uncertainty to weigh losses for scene geometry and semantics. in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* 7482–7491, IEEE/CVF (2018).
 - [45] Murakami, T., Hashidate, Y., Matsunaga, Y. Mechanistic Interpretability of Fine-Tuned Protein Language Models for Nanobody Thermostability Prediction. *bioRxiv* (2025). Preprint. <https://www.biorxiv.org/content/10.64898/2025.12.22.695878v1>