

*Regular Articles***Transfer learning from computed stability data for nanobody melting-temperature prediction**Taihei Murakami^{2,3*}, Kentaro Sasaki^{2*}, Soichiro Oda², Kazuma Okada², Yasuhiro Matsunaga^{1,2}¹ RIKEN Center for Computational Science, Kobe, Hyogo 650-0047, Japan² Saitama University, Saitama 338-8570, Japan³ Epsilon Molecular Engineering, Inc., Saitama 338-8570, Japan

Received XX Xxx, 20XX; Accepted XX Xxx, 20XX;

Released online in J-STAGE as advance publication XX Xxx, 20XX

Edited by Xxx Xxx

Thermal stability is important for nanobody production, storage, and formulation, but measuring melting temperature (T_m) for many sequences is costly. Simulations can provide stability-related data, but they do not directly calculate T_m . We tested whether simulation data related to T_m improve prediction from 57 training measurements. Multi-task models learned T_m and one computed property from ESM2 sequence representations. Settings were selected with 114 T_m validation sequences, and final models were tested on the same 396 held-out sequences. We first compared two molecular dynamics (MD) data sets based on the fraction of native contacts. With a frozen encoder, adding labels from local mutation scans of two fixed structures reduced mean absolute error by 0.20°C relative to T_m -only. Adding labels from a separately developed data set covering many structures reduced error by 0.05°C. Because the two MD studies also used different MD protocols and model searches, this comparison does not isolate the effect of variant choice. We then compared FEP, Rosetta, ThermoMPNN, and MD labels based on the fraction of native contacts. FEP gave the lowest test error with both frozen and fine-tuned encoders. With fine-tuning, mean absolute error decreased from 6.55°C to 6.40°C, although the 95% interval for the change included zero. Native-contact MD labels reduced error only with the frozen encoder; Rosetta and ThermoMPNN labels gave little or no improvement. Thus, transfer was determined by the selected variants and calculated physical quantity rather than by the number of computed labels alone, so these choices should be made before simulations begin.

Key words: nanobody thermostability, melting-temperature prediction, multi-task transfer learning, protein language model, molecular simulation labels

◀ Significance ▶

Thermal stability affects whether a nanobody can be produced, stored, and used reliably. Melting temperature is a common measure, but obtaining measurements for many sequences is costly. Simulations can add related stability data, but they cannot directly calculate melting temperature. FEP mutation-free-energy labels and MD labels based on the fraction of native contacts from local mutation scans improved T_m prediction in some model settings, whereas several other calculated data sets did not. Larger calculated data sets were not consistently more useful. The variants to simulate and the physical quantity to calculate should be chosen for the melting-temperature prediction task before simulations begin.

*These authors contributed equally.

Corresponding author: Yasuhiro Matsunaga, RIKEN Center for Computational Science, Kobe, Hyogo 650-0047, Japan.
ORCID iD: <https://orcid.org/0000-0003-2872-3908>, e-mail: ymatsunaga@riken.jp

Introduction

Nanobodies are single-domain VHH antibody fragments used as small, modular binding proteins [1, 2]. Their thermal stability affects how reliably they can be expressed, purified, and stored [3–5], yet both direct measurements and detailed stability simulations remain costly [6–9]. Melting temperature (T_m) is a common experimental measure of thermal stability, but only a limited number of nanobody T_m values are available for model training.

Protein language models such as ESM2 learn sequence features from large protein databases [10–12], and several methods use these features to predict T_m [13–18]. For family-specific prediction, measured T_m values are still needed for training and evaluation. Large experimental stability sets exist for other protein classes: Tsuboyama et al., for example, reported on the order of 10^5 high-quality folding-stability measurements for single mutants and selected double mutants of several hundred natural and designed protein domains [19]. Those domains are 40–72 residues long and differ from VHH domains in both sequence and measurement type, and a model fine-tuned on these short domains was less accurate on longer proteins [20]. Nanobody-specific data and validation therefore remain necessary [21–23].

Computed properties offer another way to supplement scarce experimental data. Combining a large in silico data set with a smaller experimental one is often called simulation-to-real (Sim2Real) learning, and this approach has improved materials-property prediction in several settings [24–26]. More computed examples are not always a better experimental predictor, however, because a calculation may be inaccurate for the systems of interest, may cover systems unrelated to the target, or may report a quantity that is only weakly connected to the experiment.

The last point is the acute one for protein stability, because what a calculation returns is rarely what an experiment measures. T_m is the midpoint of a thermal-unfolding transition under specified assay conditions. Stability calculations instead usually report the change in folding free energy caused by a mutation ($\Delta\Delta G$): free-energy perturbation (FEP) estimates $\Delta\Delta G$ from alchemical simulations [27], Rosetta estimates it from an empirical energy function [28, 29], and learned models such as ThermoMPNN predict it from structure [30]. Molecular dynamics (MD) reports a different kind of quantity altogether, such as the fraction of native contacts retained during a high-temperature trajectory [31]. Each of these concerns protein stability, but none of them is T_m , so their value as auxiliary data for T_m prediction cannot be assumed.

Two decisions on study design therefore precede any Sim2Real model training, and both must be made before the simulations begin. The first is which sequence variants to compute. A heterogeneous set drawn from many structures may carry unwanted differences, such as sequence length, whereas a mutation scan on one or two fixed structures gives a narrower and more controlled set of sequence changes. The second is which quantity to compute for those variants, since $\Delta\Delta G$ from FEP, a Rosetta score, and an MD native-contact value need not carry the same information about T_m . Which of these decisions governs whether computed stability data transfer to a measured T_m , and does computing more examples help on its own? This is the question we address here.

To address this question, we evaluated several computed quantities in the same low-data T_m task: mutation free energies from FEP, Rosetta mutation scores, ThermoMPNN scores, Rosetta scores for ESM2-proposed or random variants, and MD native-contact values. For MD, we examined a heterogeneous structure panel and local mutation scans on the same 1MEL and 4IDL structures used for FEP. Both MD data sets used the same force field, water model, simulation engine, and 400 K protocol, and both recorded the same native-contact quantity; they differed in the variants covered and in the contact selection and averaging window used to compute it. We therefore compare the two MD data sets as they were built. For each data source, the model was trained on T_m and one computed quantity using ESM2 sequence representations (Fig. 1). Model settings were selected using only experimental T_m validation data. The selected models were then compared using paired absolute errors for the same 396 test sequences, which were used neither for training nor for model selection. We use “transfer” to mean a reduction in test T_m error after adding computed data.

We found that both design decisions, and not the amount of computed data, determined whether transfer occurred. Among the computed quantities, FEP $\Delta\Delta G$ gave the lowest test error with both encoder settings, and MD native-contact labels reduced error with the frozen encoder, whereas Rosetta and ThermoMPNN labels gave little or no improvement. Among the two ways of choosing variants, the local mutation scans gave a larger frozen-encoder gain than the heterogeneous panel. The number of computed values alone did not predict transfer performance: performance instead differed among data sets that varied in their selected sequences and computed quantities. Because these choices are fixed once the simulations are run, and the simulations are the expensive step, this finding suggests spending a computational budget on a small, targeted mutation scan of a quantity close to the measured property rather than on a larger but loosely related data set. That reasoning is not specific to nanobody T_m , and we expect it to apply to other Sim2Real settings in which the computed and measured quantities differ.

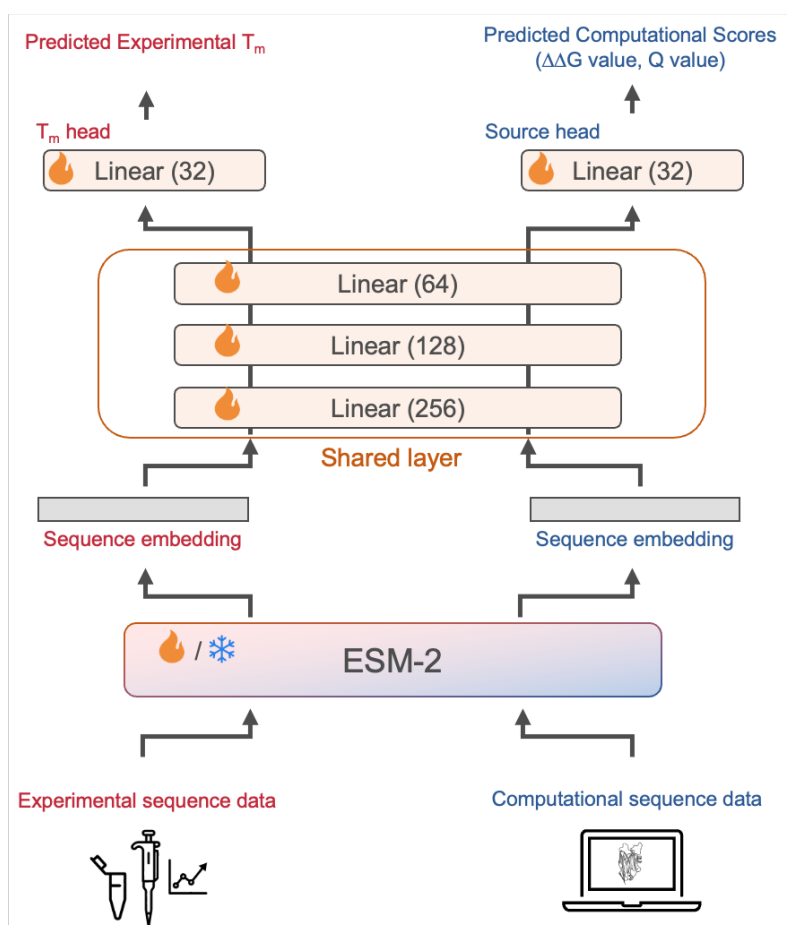


Figure 1 Conceptual multi-task setup used to evaluate calculated stability values. All models use ESM2 sequence representations and separate outputs for T_m and the calculated quantity. The diagram shows the shared-trunk architecture; the architectures tested differed in how much of the downstream model was shared. ESM2 is either kept frozen or fine-tuned. Models were selected with experimental T_m validation data and evaluated with held-out T_m data.

Materials and methods

Study design

We tested whether computed protein-stability labels could improve nanobody melting-temperature (T_m) prediction in a low-data setting. Checkpoints and hyperparameters were selected solely by error on the experimental T_m validation set. A computed-label source supplied either one auxiliary task or separate tasks for the 1MEL and 4IDL tables. The T_m and auxiliary examples used the same ESM2 encoder. Each model was compared with a T_m-only model whose encoder was frozen or fine-tuned in the same way.

Experimental T_m data

We used the thermo-tm data set from NbBench, a public nanobody benchmark [22]. To create the low-data setting, we used the published validation split (57 sequences) for training, the published test split (114 sequences) for model selection, and the published training split (396 sequences) as the final test set. Model parameters were fitted on the 57-sequence set, checkpoints and hyperparameters were selected on the 114-sequence set, and the 396-sequence set was used only for final evaluation. T_m values were stored in degrees Celsius. We fit a robust scaler and then a min-max scaler to [0,1] using only the 57 training values. The fitted scalers were applied to the training and validation labels. Predictions were returned to degrees Celsius before validation and test errors were calculated. Validation and test T_m values therefore did not affect the scaling step.

Computed-label tables and sampling

We tested seven sources of computed labels: FEP mutation free energies, matched MD native-contact values, Rosetta mutation scores, ThermoMPNN predictions, Rosetta scores for random variants, Rosetta scores for ESM2-proposed variants, and MD native-contact values for a heterogeneous structure panel. We used these values as auxiliary targets, not as estimates of T_m . The FEP, Rosetta, and ThermoMPNN sources each comprised two single-mutation tables: 435 rows for 1MEL and 409 rows for 4IDL. The matched MD tables contained 431 and 406 rows, respectively, after restricting the analysis to variants that reached the common analysis time. All 837 matched-MD sequences occurred in the 844-row FEP pool; seven FEP variants lacked a matched MD label. The random and ESM2-proposed sets each contained 1000 two-mutation variants per structure. The heterogeneous MD table contained 1143 PDB-derived rows representing 833 unique amino-acid sequences. Its sampling unit was a PDB row, so sequences represented by multiple structures could occur more than once.

The 1MEL and 4IDL tables were treated as separate tasks because their score ranges can differ. Models used either a separate output head for each structure or one head conditioned on a learned structure indicator, as described below. Each processed table supplied an amino-acid sequence and one scalar label scaled to $[0,1]$. FEP, Rosetta, and ThermoMPNN scores were sign-reversed so that larger values indicated a more favorable change. Each structure table was then processed separately by robust scaling, a Yeo–Johnson transform with standardization, and min–max scaling. MD contact values were min–max scaled, with larger values indicating greater contact retention.

For a source with two structure tables, n denotes the maximum number of rows sampled from *each* table for each ensemble member, not the total number. At $n = 320$, for example, one member sampled at most 640 computed-label rows. For the single-table heterogeneous source, n is the total number sampled per member. Each run used its own seed, and rows were sampled independently for each label source. The FEP and matched-MD runs therefore did not necessarily use the same mutations from their nearly identical pools. Of the sampled rows, 80% were used for training and 20% were held out. The held-out auxiliary rows were not used for checkpoint or hyperparameter selection; those decisions used only the experimental T_m validation set.

We also checked exact sequence matches between every computed-label table and the NbBench splits. There were no exact matches for the FEP, matched MD, Rosetta, ThermoMPNN, random-variant, or ESM2-proposed tables. The heterogeneous MD table contained 2 training, 4 validation, and 8 test sequences also present in NbBench. When one of these rows was sampled for auxiliary training, the model saw the target sequence with a computed Q label. It did not see the experimental validation or test T_m value, which remained unused in training. We retained these rows; the heterogeneous panel was therefore not fully disjoint from the target sequences. The relation between sequence length and this MD label is shown in Supplementary Fig. S1.

FEP mutation free energies

The FEP labels were mutation free-energy changes, $\Delta\Delta G$, for the 1MEL and 4IDL mutation scans. Positive raw $\Delta\Delta G$ denotes a destabilizing mutation. The FEP tables were produced with alchemical simulations in NAMD using the CHARMM36 protein force field and explicit TIP3P water [32–34]. The starting coordinates were the 1MEL and 4IDL crystal structures. Mutation systems were prepared with the VMD AlaScan workflow and CHARMM hybrid-residue topologies [35, 36]. The simulations used a 2-fs timestep, particle mesh Ewald electrostatics, fixed bonds to hydrogen, a temperature of 300 K, and a pressure of 1 atm [37, 38]. Each forward and backward transformation covered 20 windows of width $\Delta\lambda = 0.05$ between 0 and 1. A window contained 500,000 steps (1 ns), with its first 100,000 steps assigned to equilibration. The folded-state free energy was combined with a residue-specific unfolded-state reference to obtain $\Delta\Delta G$ relative to the wild type.

MD native-contact labels

Both MD labels were based on the smooth native-contact fraction Q [31]. For a native contact with reference distance d_{ij}^0 and distance $d_{ij}(t)$ at time t , its contribution was

$$q_{ij}(t) = \left[1 + \exp \left\{ \beta [d_{ij}(t) - \lambda d_{ij}^0] \right\} \right]^{-1},$$

with $\beta = 50 \text{ nm}^{-1}$ and $\lambda = 1.8$. Native contacts were pairs within 0.45 nm in the reference frame and separated by more than three residues. Q was the mean over the selected contacts and frames.

Matched mutation scans. We simulated the 1MEL and 4IDL wild types and the same single-mutant sequences used in the FEP comparison. Each protein was placed in a cubic box of OPC water with a 15-Å buffer and 0.15 M salt,

and was described by Amber ff19SB with hydrogen-mass repartitioning [39,40]. OpenMM was used for minimization, equilibration, and production [41]. We equilibrated each system for 0.2 ns under NVT conditions and 0.8 ns under NPT conditions at 300 K, followed by 0.5 ns of restrained NVT dynamics at 400 K. We then ran unrestrained NVT production at 400 K. All dynamics stages used a 4-fs timestep, and production coordinates were saved every 10 ps. Although some runs continued to 100 ns, every label was calculated from the first 4000 production frames (40 ns). For this source, the reference distances came from production frame 0 and all backbone atoms selected by MDAnalysis were eligible for contacts. We computed $\Delta Q = Q_{\text{mutant}} - Q_{\text{WT}}$ within each structure and then min–max scaled it. Subtracting one constant wild-type value does not change a min–max-scaled label, so scaled Q and scaled ΔQ are identical within each structure table.

Heterogeneous nanobody panel. The comparison panel was built from SABDab structures [42]. For each PDB entry with an IMGT-renumbered structure, we selected the first heavy-chain entry whose chain identifier was present in the structure file. Preparation used PDBFixer and first attempted protonation at pH 7.4 with pdb2pqr and PROPKA [43,44]. If that step failed or was unavailable, the workflow used pdb4amber with Reduce; if Amber conversion also failed, it retained the PDBFixer output. The systems used ff19SB, OPC water, 0.15 M salt, and OpenMM. After 300 K NVT and NPT equilibration, two restrained 400 K NVT stages preceded unrestrained 400 K NVT production scheduled for 100 ns. Requiring a usable topology and trajectory, a sequence of at least 50 residues, and at least one selected contact gave 1143 rows. The first production frame supplied the reference distances. The contact set contained backbone heavy-atom pairs for which at least one residue was Asp, Glu, Gln, Asn, Arg, Lys, or His. For trajectories longer than 300 frames, Q was averaged over the larger of the final 300 frames and the final one third of the trajectory; all frames were used for shorter trajectories. The label was min–max scaled across the single table. The number of frames entering a label ranged from 66 to 333 (median 333); nine rows had fewer than 300 available frames. Both MD data sets used the Best–Hummer form at 400 K. They differed in contact selection, trajectory length, and the frames used for averaging.

Rosetta, ThermoMPNN, and variant-selection comparisons

Rosetta scores were calculated with the `ddg_monomer` application and the `ref2015` score function [28,29]. We ran 50 iterations with Cartesian minimization and full-atom input and retained the minimum score. The same settings were used for the 1MEL/4IDL mutation tables and for both sets of two-mutation variants. For the random set, mutation positions and replacement amino acids were sampled uniformly, excluding cysteine as a replacement, until 1000 unique two-mutation variants had been obtained for each structure. For the ESM2-guided set, the 650M-parameter ESM2 model generated 100,000 two-mutation variants per structure [11]. Variants were ranked by pseudo-log-likelihood, and the top 1% were scored with Rosetta. ThermoMPNN predictions were made for the same 435 and 409 single-mutation sequences used by the structure-matched tables [30]. We used the processed 435-row 1MEL and 409-row 4IDL tables.

Sequence model and regression heads

The main encoder was the 8M-parameter ESM2 model `facebook/esm2_t6.8M_UR50D` [11]. The tokenizer used a maximum length of 160 tokens, including the beginning and end tokens, so at most 158 amino acids were supplied to ESM2. Longer sequences were truncated at the end. All NbBench target sequences were 151 residues or shorter. The first-token representation was used as the sequence vector. In the shared model, this vector passed through linear layers of size 256, 128, and 32, with ReLU activation and dropout after the first two layers. Linear output layers produced the T_m and auxiliary predictions; the two-table auxiliary-head designs are described below. We also tested residual and latent architectures, each with two 256–128–32 multilayer perceptrons (MLPs) applied to the same ESM2 representation. The auxiliary head used the second MLP. The residual architecture added a sigmoid-gated output from that MLP to a base T_m prediction. The latent architecture predicted a T_m correction from both MLP outputs and their elementwise product. For two-table sources, we compared separate 1MEL and 4IDL heads with a single head conditioned on a learned structure embedding. In the residual and latent architectures, gradients from the auxiliary loss were stopped before the ESM2 encoder by default, whereas the T_m loss could still update the encoder. In the shared architecture, all task losses could update ESM2 when it was fine-tuned.

We tested two encoder settings. In the frozen setting, all ESM2 parameters were fixed and sequence vectors were computed once before head training. In the fine-tuned setting, all ESM2 parameters could change. AdamW used a smaller learning rate for ESM2 than for the newly initialized layers. Selected controls used 35M- and 650M-parameter ESM2 encoders without size-specific tuning; their settings are given in the Supplementary methods.

Loss and model selection

Each row had one task identifier and contributed loss only through that task's output head. We used Huber loss with $\delta = 1$ on the scaled labels. By default, the losses were combined with learned uncertainty weights [45]. For task k , loss L_k , and learned log-scale s_k , the term added to the total loss was

$$\frac{L_k}{2 \exp(2s_k)} + s_k.$$

Fixed task weights were tested as controls. Models were trained with AdamW, cosine learning-rate decay, a batch size of 16, 100 warmup steps, and at most 400 epochs. Mixed-precision training was used on CUDA devices. A checkpoint was saved and evaluated once per epoch. The checkpoint with the lowest mean squared error on the 114 experimental Tm validation rows was retained, and training stopped after 15 epochs without improvement. The run seed controlled randomness in Python, NumPy, PyTorch, and CUDA, as well as computed-row sampling and the auxiliary training–holdout split.

For the structure-matched label comparison in Fig. 3, each label source was tuned separately in the frozen and fine-tuned settings. The search had two stages. Stage 1 considered three architectures (shared, residual, and latent) and two auxiliary-head designs (separate and structure-conditioned); Tm-only models considered the same three architectures. Stage 2 held the selected architecture and head design fixed. Candidate fine-tuned models used encoder learning rates of 1×10^{-4} , 5×10^{-5} , 3×10^{-5} , and 1×10^{-5} . Candidate frozen models used head learning rates of 1×10^{-3} , 5×10^{-4} , and 2×10^{-4} . We also varied dropout among 0.05, 0.15, and 0.30 and weight decay between 0.02 and 0.10. Not every combination of these settings was run for every source. We ranked the completed candidates by the MAE of a three-seed ensemble on the experimental Tm validation set, using $n = 320$ rows per structure table. The selected setting was then trained with five seeds for the final test result. Label-count curves used the same selected setting at $n = 20, 80, 160$, and 320 per structure; they were not retuned at each point.

The heterogeneous-panel study used a separate three-seed validation search at $n = 640$. The MD model used the residual architecture. We varied the non-encoder and encoder learning rates, dropout, and auxiliary-loss weighting one factor at a time. Tm-only references were selected separately from shared, residual, and latent architectures. The selected frozen MD model used a non-encoder learning rate of 3×10^{-4} , dropout 0.30, and weight decay 0.04. The selected fine-tuned MD model used non-encoder and encoder learning rates of 1×10^{-4} and 3×10^{-5} , respectively, dropout 0.15, and weight decay 0.04. Both used learned uncertainty weighting. The selected Tm-only references used the shared architecture, a non-encoder learning rate of 3×10^{-4} , weight decay 0.04, and dropout 0.30 (frozen) or 0.05 (fine-tuned); the fine-tuned reference used an encoder learning rate of 1×10^{-4} . The selected configurations were retrained with ten seeds for the final results. Sampling $n = 640$ rows from the single heterogeneous table matched the maximum total sampled per ensemble member at $n = 320$ from each of the two matched tables.

Test evaluation and statistics

For each main 8M condition, every trained model predicted all 396 test sequences. We averaged predictions across five training seeds for the structure-matched sources and ten for the heterogeneous panel. The averaged predictions were returned from scaled values to degrees Celsius. The main measure was mean absolute error (MAE). We also calculated root mean squared error, R^2 , and Spearman and Pearson correlations. Uncertainty was estimated by nonparametric bootstrap over test proteins. For one condition, we resampled its 396 absolute errors with replacement 10,000 times, using random seed 42, and calculated MAE for each sample. The reported 95% confidence interval is the 2.5th to 97.5th percentile range. For a comparison, the same resampled test indices were applied to the candidate and reference error vectors. We calculated Δ MAE as candidate MAE minus reference MAE for every sample. Negative values therefore favor the candidate. Main figures report the mean paired Δ MAE and its 95% confidence interval.

Software

Model training and evaluation were implemented in Python with PyTorch and Hugging Face Transformers. The released workflow starts from the processed label tables and reproduces model training, evaluation, and figure generation; it does not repeat the underlying FEP, MD, Rosetta, or ThermoMPNN calculations.

Results and discussion

Models were selected using experimental Tm error

All comparisons used experimental Tm as the target and the same 396-sequence test set; computed labels entered only as auxiliary training targets (Fig. 1). For each data source, model settings were selected with the experimental Tm validation set. The Tm-only test MAE was 7.23°C with frozen ESM2 and 6.55°C with fine-tuned ESM2 in the structure-matched comparison. The separately selected Tm-only references for the heterogeneous-panel study had MAEs of 7.38 and 6.61°C, respectively. Each computed-label model was compared only with the Tm-only model selected in the same study and encoder setting.

The frozen-encoder MAE reduction was numerically larger for matched scans

We compared two ways of building an MD label set. The heterogeneous data set contained 1143 PDB-derived rows representing 833 unique sequences of different lengths. The matched data set contained local single mutations of the fixed 1MEL and 4IDL structures. Both used 400 K MD and the Best-Hummer native-contact Q form, but their contact definitions, trajectory lengths, averaging windows, and model searches differed. The comparison therefore does not isolate the effect of sequence selection. Because the two studies were developed and selected separately, we compared the change from the Tm-only model within each study rather than their absolute MAEs (Fig. 2a). In the heterogeneous study, adding the MD label changed MAE by -0.049°C with a frozen encoder (95% CI, -0.092 to -0.006°C) and by $+0.116^\circ\text{C}$ with a fine-tuned encoder (95% CI, -0.009 to $+0.243^\circ\text{C}$). In the matched study, the corresponding changes were -0.195°C (95% CI, -0.366 to -0.030°C) and $+0.029^\circ\text{C}$ (95% CI, -0.139 to $+0.197^\circ\text{C}$). The frozen-encoder MAE reduction was therefore larger for the matched mutation scans.

Eight sequences in the heterogeneous MD table also occurred in the Tm test set. Excluding their test errors barely changed either estimate: the paired changes were -0.049°C for the frozen encoder (95% CI, -0.094 to -0.005°C) and $+0.122^\circ\text{C}$ for the fine-tuned encoder (95% CI, -0.005 to $+0.248^\circ\text{C}$). We did not retrain after removing the overlapping computed-label rows, so this check addresses their contribution to test MAE but not a possible effect during training or model selection. Other differences between the data sets may also matter. In the heterogeneous table, Q had a Pearson correlation of $r = +0.13$ with sequence length (Supplementary Fig. S1). The matched scans held length fixed within each structure and described the local effect of one mutation. The heterogeneous table also sampled PDB rows, giving more weight to sequences represented by multiple structures. These observations do not identify the cause of the difference.

The frozen-encoder label-count curves provide a second comparison of FEP and matched MD (Fig. 2b). At $n = 20$, the paired MAE changes were -0.10°C for FEP and $+0.09^\circ\text{C}$ for matched MD. At $n = 320$, they were -0.22°C and -0.20°C , respectively, and both 95% intervals were below zero. The direct FEP-minus-MD difference at $n = 320$ was -0.03°C (95% CI, -0.24 to $+0.19^\circ\text{C}$), which did not resolve a difference between the labels. Four label counts, with nonmonotonic changes, are too few to establish a scaling relation.

FEP had the lowest test MAE with both encoders

We next compared all computed labels after selecting the model for each source separately (Fig. 3a). With a frozen encoder, FEP reached 7.01°C MAE, a paired change of -0.221°C from Tm-only (95% CI, -0.393 to -0.051°C). Matched MD reached 7.03°C, a change of -0.195°C (95% CI, -0.366 to -0.030°C). Both intervals were below zero. ThermoMPNN reached 7.09°C, but its interval included zero. The Rosetta-based labels were close to or worse than the Tm-only model. With a fine-tuned encoder, FEP reached 6.40°C, the lowest MAE in the study. Its paired change from the 6.55°C Tm-only model was -0.153°C (95% CI, -0.323 to $+0.020^\circ\text{C}$), an interval that included zero. Matched MD, ThermoMPNN, plain Rosetta, and random variants scored with Rosetta differed from Tm-only by $+0.029$ to $+0.144^\circ\text{C}$; their intervals also included zero. ESM2-proposed variants scored with Rosetta increased error by $+0.411^\circ\text{C}$ (95% CI, $+0.140$ to $+0.698^\circ\text{C}$).

The direct FEP-minus-MD comparison is shown in Fig. 3b. Negative values favor FEP. The difference was -0.026°C with a frozen encoder (95% CI, -0.243 to $+0.192^\circ\text{C}$) and -0.182°C with a fine-tuned encoder (95% CI, -0.393 to $+0.024^\circ\text{C}$). Neither interval excluded zero at the largest label count. The fine-tuned label-count curves show how the two results changed with the amount of auxiliary data (Fig. 3c). At $n = 20$, matched MD increased MAE by $+0.47^\circ\text{C}$ (95% CI, $+0.26$ to $+0.69^\circ\text{C}$). Its error decreased as n increased, and the interval included zero by $n = 160$. FEP changed from $+0.02^\circ\text{C}$ at $n = 20$ to -0.15°C at $n = 320$, but the interval at the largest setting also included zero. These four label counts show the direction of the changes but are too few to establish a scaling relation.

FEP $\Delta\Delta G$ measures the free-energy effect of a mutation, whereas high-temperature Q measures contact retention over a finite trajectory. The closer thermodynamic connection between $\Delta\Delta G$ and mutation stability may help explain why FEP had a lower MAE after ESM2 fine-tuning, although the present experiments do not test that mechanism. In a recent




Figure 2 The frozen-encoder MAE reduction was larger for the matched mutation scans than for the heterogeneous MD panel. The heterogeneous panel covered PDB structures of different lengths, whereas the matched scans used local mutations of 1MEL and 4IDL. Both used 400 K MD and the same Q form, but other parts of the protocols and model searches differed. (a) Paired MAE change from the Tm-only model within each study. Zero means equal MAE, and negative values mean lower Tm error after adding the MD label. Predictions were averaged across ten training seeds for the heterogeneous panel and five for the matched scans. (b) Paired MAE change from the frozen Tm-only model at four label counts. Here n is the number sampled from each structure table for each ensemble member; 80% of the sampled rows entered training. In both panels, points and whiskers show the paired change and its 95% bootstrap interval over the 396 test sequences.

preprint, we used sparse autoencoders to examine ESM2 after fine-tuning on nanobody Tm and found features associated with surface charge, hydrophobic-core packing, the VHH tetrad, and disulfide bonds [46]. We did not examine the representations learned here. A useful follow-up would be to test whether FEP supervision strengthens the same features by comparing Tm-only, Tm+FEP, and Tm+MD models with the same regression architecture.

Additional checks and limitations

Without size-specific tuning, FEP reduced MAE relative to Tm-only for both the 35M- and 650M-parameter ESM2 models; both paired 95% intervals were below zero (Supplementary Fig. S2a). The larger models used fixed settings rather than the search applied to the 8M model, so these tests do not compare performance across model sizes. An analytical correction for periodic images in charge-changing FEP mutations had almost no effect on the scaled labels (Supplementary Fig. S2b). Under the Rosetta protocol used here, ESM2-proposed variants did not give lower Tm error than random variants. This comparison provides no evidence that the ESM2 proposal step improved transfer to Tm. Data set size did not explain the contrast between the two MD studies. The heterogeneous pool contained 1143 rows and the matched pool contained 837, but each ensemble member sampled at most 640 auxiliary rows in both studies. The larger frozen-encoder reduction for the matched scans therefore cannot be explained by more training labels per model. FEP and matched MD required similar simulation time per variant, but wall-clock costs were not recorded consistently enough to support a general cost comparison.

The study has several limitations. The Tm training set contains only 57 nanobodies, and the matched calculations cover two structures. The two MD data sets differ in more than their selected sequences. The heterogeneous table sampled repeated PDB entries, contained eight exact matches to the Tm test sequences, included 11 rows longer than the model input, and contained trajectories that did not all reach the scheduled production length. The selected FEP and MD models also differed in regression architecture and in whether the auxiliary loss updated the encoder. The bootstrap intervals




Figure 3 FEP had the lowest test T_m MAE with both encoders. (a) Paired MAE change from the corresponding T_m-only model for each computed label. Negative values mean lower T_m error. Filled squares and open circles denote frozen and fine-tuned encoders, respectively. (b) Direct FEP-minus-MD comparison. Zero means equal T_m test-set MAE, negative values favor FEP, and positive values favor MD native-contact *Q*. (c) Paired MAE change from the fine-tuned T_m-only model at four label counts. Here *n* is the number sampled from each structure table for each ensemble member; 80% entered training. Whiskers in all panels are 95% bootstrap intervals over the 396 test sequences.

resample the fixed test proteins after predictions are averaged across trained models. They do not include uncertainty from selecting auxiliary rows, training seeds, or model settings. Finally, we did not make new T_m measurements or evaluate predictions on newly measured nanobodies.

Conclusion

Computed labels did not improve low-data nanobody T_m prediction uniformly. Among two independently developed MD data sets, matched mutation scans showed a larger reduction with the frozen encoder than the heterogeneous structure panel, although differences in their protocols prevent attribution to sequence selection alone. FEP had the lowest test MAE with both frozen and fine-tuned ESM2; its 95% interval excluded zero only with the frozen encoder. The comparisons did not show a simple relation between label count and T_m error. When calculations are intended to support T_m prediction, both the variants and the physical property should be considered from the start.

Matched scans could next be extended to more nanobody structures, with new calculations selected adaptively. Reinforcement learning could support this sequential selection, but it should be compared with random sampling and simpler active-learning methods. Because ESM2 proposals did not yield lower T_m prediction error than random variants under the present Rosetta protocol, a computed score alone should not define success. Any improvement should ultimately be tested with new experimental T_m measurements.

Conflict of interest

T.M. is an employee of Epsilon Molecular Engineering, Inc. Y.M. is a scientific advisor of Epsilon Molecular Engineering, Inc.

Author contributions

Y.M. conceived and designed the study. Y.M., T.M., and K.S. performed the calculations. Y.M. and T.M. analyzed the data. S.O. and K.O. performed the FEP calculations. Y.M. and T.M. wrote the manuscript. All authors discussed the results, reviewed the manuscript, and approved the final version.

Data availability

The source code for data processing, model training, analysis, and figure generation is available at <https://github.com/matsunagalab/sim2real>. The processed label tables, per-sequence test errors, and figure source data will be available from Zenodo (DOI: 10.5281/zenodo.XXXXXXX). Raw MD and FEP files are not included because of their size and are available from the corresponding author upon reasonable request.

Acknowledgements

This work was supported by JSPS KAKENHI (23H03412), JST NBDC (JPMJND2603), and MEXT's "Program for Promoting Researches on the Supercomputer Fugaku" (Development and application of large-scale simulation-based inferences for biomolecules; JPMXP1020230119). Computational resources were provided by the supercomputer Fugaku at the RIKEN Center for Computational Science under Project IDs hp230209, hp240215, and hp250233.

References

- [1] Muyldermans, S. Nanobodies: Natural single-domain antibodies. *Annual Review of Biochemistry* 82, 775–797 (2013). <https://doi.org/10.1146/annurev-biochem-063011-092449>
- [2] Alexander, E. Leong, K. W. Discovery of nanobodies: a comprehensive review of their applications and potential over the past five years. *Journal of Nanobiotechnology* 22, 661 (2024). <https://doi.org/10.1186/s12951-024-02900-y>
- [3] Kunz, P., Zinner, K., Mücke, N., Bartoschik, T., Muyldermans, S., Hoheisel, J. D. The structural basis of nanobody unfolding reversibility and thermoresistance. *Sci. Rep.* 8, 7934 (2018).
- [4] Mezzasalma, T. M., Kranz, J. K., Chan, W., Struble, G. T., Schalk-Hihi, C., Deckman, I. C., et al. Enhancing recombinant protein quality and yield by protein stability profiling. *J. Biomol. Screen.* 12, 418–428 (2007).
- [5] Kumar, S., Tsai, C. J., Nussinov, R. Factors enhancing protein thermostability. *Protein Eng.* 13, 179–191 (2000).
- [6] Bekker, G.-J., Ma, B., Kamiya, N. Thermal stability of single-domain antibodies estimated by molecular dynamics simulations. *Protein Science* 28, 429–438 (2019). <https://doi.org/10.1002/pro.3546>
- [7] Higashida, R. Matsunaga, Y. Enhanced conformational sampling of nanobody CDR H3 loop by generalized replica-exchange with solute tempering. *Life (Basel)* 11, 1428 (2021).

- [8] Mohseni, A., Molakarimi, M., Taghdir, M., Sajedi, R. H., Hasannia, S. Exploring single-domain antibody thermostability by molecular dynamics simulation. *J. Biomol. Struct. Dyn.* 37, 3686–3696 (2019).
- [9] Tomimoto, Y., Yamazaki, R., Shirai, H. Increasing the melting temperature of VHH with the in silico free energy score. *Sci. Rep.* 13, 4922 (2023).
- [10] Rives, A., Meier, J., Sercu, T., Goyal, S., Lin, Z., Liu, J., et al. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences of the United States of America* 118, e2016239118 (2021). <https://doi.org/10.1073/pnas.2016239118>
- [11] Lin, Z., Akin, H., Rao, R., Hie, B., Zhu, Z., Lu, W., et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science* 379, 1123–1130 (2023). <https://doi.org/10.1126/science.ade2574>
- [12] Hayes, T., Rao, R., Akin, H., Sofroniew, N. J., Oktay, D., Lin, Z., et al. Simulating 500 million years of evolution with a language model. *Science* 387, 850–858 (2025).
- [13] Blaabjerg, L. M., Kassem, M. M., Good, L. L., Jonsson, N., Cagiada, M., Johansson, K. E., et al. Rapid protein stability prediction using deep learning representations. *eLife* 12, e82593 (2023). <https://doi.org/10.7554/eLife.82593>
- [14] Li, M., Wang, H., Yang, Z., Zhang, L., Zhu, Y. DeepTM: A deep learning algorithm for prediction of melting temperature of thermophilic proteins directly from sequences. *Computational and Structural Biotechnology Journal* 21, 5544–5560 (2023). <https://doi.org/10.1016/j.csbj.2023.11.006>
- [15] Alvarez, J. A. E. Dean, S. N. TEMPRO: nanobody melting temperature estimation model using protein embeddings. *Scientific Reports* 14, 19074 (2024). <https://doi.org/10.1038/s41598-024-70101-6>
- [16] Savojardo, C., Manfredi, M., Martelli, P. L., Casadio, R. DDGemb: predicting protein stability change upon single- and multi-point variations with embeddings and deep learning. *Bioinformatics* 41 (2024).
- [17] Ramon, A., Ni, M., Predeina, O., Gaffey, R., Kunz, P., Onuoha, S., et al. Prediction of protein biophysical traits from limited data: a case study on nanobody thermostability through NanoMelt. *MAbs* 17, 2442750 (2025).
- [18] Mao, J., Song, Y., Kong, M., Guo, Y., Liu, Y., Pu, X. Thermostability prediction powered by synergistic deep learning at experimental and theoretical levels for nanobodies. *ACS Appl. Mater. Interfaces* 18, 6433–6451 (2026).
- [19] Tsuboyama, K., Dauparas, J., Chen, J., Laine, E., Mohseni Behbahani, Y., Weinstein, J. J., et al. Mega-scale experimental analysis of protein folding stability in biology and design. *Nature* 620, 434–444 (2023). <https://doi.org/10.1038/s41586-023-06328-6>
- [20] Chu, S. K. S., Narang, K., Siegel, J. B. Protein stability prediction by fine-tuning a protein language model on a mega-scale dataset. *PLOS Computational Biology* 20, e1012248 (2024). <https://doi.org/10.1371/journal.pcbi.1012248>
- [21] Valdés-Tresanco, M. S., Valdés-Tresanco, M. E., Molina-Abad, E., Moreno, E. NbThermo: a new thermostability database for nanobodies. *Database* 2023, baad021 (2023). <https://doi.org/10.1093/database/baad021>
- [22] Zhang, Y., Tsuda, K. NbBench: Benchmarking language models for comprehensive nanobody tasks. *arXiv* 2505.02022 (2025). <https://doi.org/10.48550/arXiv.2505.02022>
- [23] Schmirler, R., Heinzinger, M., Rost, B. Fine-tuning protein language models boosts predictions across diverse tasks. *Nature Communications* 15, 7407 (2024). <https://doi.org/10.1038/s41467-024-51844-2>
- [24] Yamada, H., Liu, C., Wu, S., Koyama, Y., Ju, S., Shiomi, J., et al. Predicting materials properties with little data using shotgun transfer learning. *ACS Central Science* 5, 1717–1730 (2019). <https://doi.org/10.1021/acscentsci.9b00804>
- [25] Aoki, Y., Wu, S., Tsurimoto, T., Hayashi, Y., Minami, S., Okubo, T., et al. Multitask machine learning to predict polymer–solvent miscibility using Flory–Huggins interaction parameters. *Macromolecules* 56, 5446–5456 (2023). <https://doi.org/10.1021/acs.macromol.2c02600>
- [26] Minami, S., Hayashi, Y., Wu, S., Fukumizu, K., Sugisawa, H., Ishii, M., et al. Scaling law of Sim2Real transfer learning in expanding computational materials databases for real-world predictions. *npj Computational Materials* 11, 146 (2025). <https://doi.org/10.1038/s41524-025-01606-5>
- [27] Bennett, C. H. Efficient estimation of free energy differences from Monte Carlo data. *Journal of Computational Physics* 22, 245–268 (1976). [https://doi.org/10.1016/0021-9991\(76\)90078-4](https://doi.org/10.1016/0021-9991(76)90078-4)
- [28] Kellogg, E. H., Leaver-Fay, A., Baker, D. Role of conformational sampling in computing mutation-induced changes in protein structure and stability. *Proteins: Structure, Function, and Bioinformatics* 79, 830–838 (2011). <https://doi.org/10.1002/prot.22921>
- [29] Alford, R. F., Leaver-Fay, A., Jeliakov, J. R., O’Meara, M. J., DiMaio, F. P., Park, H., et al. The Rosetta all-atom energy function for macromolecular modeling and design. *Journal of Chemical Theory and Computation* 13, 3031–3048 (2017). <https://doi.org/10.1021/acs.jctc.7b00125>
- [30] Dieckhaus, H., Brocidiaco, M., Randolph, N. Z., Kuhlman, B. Transfer learning to leverage larger datasets for improved prediction of protein stability changes. *Proceedings of the National Academy of Sciences of the United States of America* 121, e2314853121 (2024). <https://doi.org/10.1073/pnas.2314853121>

- [31] Best, R. B., Hummer, G., Eaton, W. A. Native contacts determine protein folding mechanisms in atomistic simulations. *Proceedings of the National Academy of Sciences of the United States of America* 110, 17874–17879 (2013). <https://doi.org/10.1073/pnas.1311599110>
- [32] Phillips, J. C., Hardy, D. J., Maia, J. D. C., Stone, J. E., Ribeiro, J. V., Bernardi, R. C., et al. Scalable molecular dynamics on CPU and GPU architectures with NAMD. *The Journal of Chemical Physics* 153, 044130 (2020). <https://doi.org/10.1063/5.0014475>
- [33] Best, R. B., Zhu, X., Shim, J., Lopes, P. E. M., Mittal, J., Feig, M., et al. Optimization of the additive CHARMM all-atom protein force field targeting improved sampling of the backbone ϕ , ψ and side-chain χ_1 and χ_2 dihedral angles. *Journal of Chemical Theory and Computation* 8, 3257–3273 (2012). <https://doi.org/10.1021/ct300400x>
- [34] Jorgensen, W. L., Chandrasekhar, J., Madura, J. D., Impey, R. W., Klein, M. L. Comparison of simple potential functions for simulating liquid water. *The Journal of Chemical Physics* 79, 926–935 (1983). <https://doi.org/10.1063/1.445869>
- [35] Humphrey, W., Dalke, A., Schulten, K. VMD: Visual molecular dynamics. *Journal of Molecular Graphics* 14, 33–38 (1996). [https://doi.org/10.1016/0263-7855\(96\)00018-5](https://doi.org/10.1016/0263-7855(96)00018-5)
- [36] Ramadoss, V., Dehez, F., Chipot, C. AlaScan: A graphical user interface for alanine scanning free-energy calculations. *Journal of Chemical Information and Modeling* 56, 1122–1126 (2016). <https://doi.org/10.1021/acs.jcim.6b00162>
- [37] Essmann, U., Perera, L., Berkowitz, M. L., Darden, T., Lee, H., Pedersen, L. G. A smooth particle mesh Ewald method. *The Journal of Chemical Physics* 103, 8577–8593 (1995). <https://doi.org/10.1063/1.470117>
- [38] Ryckaert, J.-P., Ciccotti, G., Berendsen, H. J. C. Numerical integration of the cartesian equations of motion of a system with constraints: molecular dynamics of n-alkanes. *Journal of Computational Physics* 23, 327–341 (1977). [https://doi.org/10.1016/0021-9991\(77\)90098-5](https://doi.org/10.1016/0021-9991(77)90098-5)
- [39] Tian, C., Kasavajhala, K., Belfon, K. A. A., Raguette, L., Huang, H., Migués, A. N., et al. ff19SB: amino-acid-specific protein backbone parameters trained against quantum mechanics energy surfaces in solution. *Journal of Chemical Theory and Computation* 16, 528–552 (2020). <https://doi.org/10.1021/acs.jctc.9b00591>
- [40] Izadi, S., Anandakrishnan, R., Onufriev, A. V. Building water models: a different approach. *The Journal of Physical Chemistry Letters* 5, 3863–3871 (2014). <https://doi.org/10.1021/jz501780a>
- [41] Eastman, P., Swails, J., Chodera, J. D., McGibbon, R. T., Zhao, Y., Beauchamp, K. A., et al. OpenMM 7: rapid development of high performance algorithms for molecular dynamics. *PLOS Computational Biology* 13, e1005659 (2017). <https://doi.org/10.1371/journal.pcbi.1005659>
- [42] Dunbar, J., Krawczyk, K., Leem, J., Baker, T., Fuchs, A., Georges, G., et al. SABDab: the structural antibody database. *Nucleic Acids Research* 42, D1140–D1146 (2014). <https://doi.org/10.1093/nar/gkt1043>
- [43] Dolinsky, T. J., Nielsen, J. E., McCammon, J. A., Baker, N. A. PDB2PQR: an automated pipeline for the setup of Poisson–Boltzmann electrostatics calculations. *Nucleic Acids Research* 32, W665–W667 (2004). <https://doi.org/10.1093/nar/gkh381>
- [44] Olsson, M. H. M., Søndergaard, C. R., Rostkowski, M., Jensen, J. H. PROPKA3: consistent treatment of internal and surface residues in empirical pK_a predictions. *Journal of Chemical Theory and Computation* 7, 525–537 (2011). <https://doi.org/10.1021/ct100578z>
- [45] Kendall, A., Gal, Y., Cipolla, R. Multi-task learning using uncertainty to weigh losses for scene geometry and semantics. in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* 7482–7491, IEEE/CVF (2018).
- [46] Murakami, T., Hashidate, Y., Matsunaga, Y. Mechanistic Interpretability of Fine-Tuned Protein Language Models for Nanobody Thermostability Prediction. *bioRxiv* (2025). Preprint. <https://www.biorxiv.org/content/10.64898/2025.12.22.695878v1>